
AdmitShift: Benchmarking Persona-Induced Drift in Emergency Admission Decisions

Priscilla Kissi-Appiah, Sean Chin, Molly Croes, and Will Tsafack

Abstract

Personalized prompting can make large language models more helpful, but in high stakes clinical allocation it can also move decisions away from a stated policy. We introduce *AdmitShift*, a benchmark for measuring persona-induced drift in emergency admission prioritization. The core design fixes the patient set and bed capacity, then changes only the user persona instruction. Each model receives 12 patient summaries from a MIMIC-IV-ED base-case table, must rank all patients, and select exactly $K = 4$ patients, where K is the bed capacity for admission. This process is defined as an episode. We have 18 disjoint episodes that are mutually exclusive among patient sets. We evaluate three local open-weight models across disjoint episodes and prompt conditions using structural reliability, admission set drift, full rank stability, acuity based clinical soundness, and fairness-oriented subgroup slices. In the primary disjoint episode experiment, Llama 3.1 8B was the most stable overall, with 72/72 valid raw rankings, the lowest bounded top- K drift (0.157), the highest Kendall rank correlation (0.769), and the lowest acuity boundary violation rate (0.074). Qwen2.5 7B was competitive on admission overlap and strongest on the severity proxy score, while Llama 3.2 3B showed the largest drift and the weakest clinical soundness profile. A secondary generated dataset stress test produced 600 responses and showed that consistency varied more by model than by persona. The generated dataset experiment adds a controlled setting in which patient properties are known, normalized after dataset generation, and used only after model responses are collected to measure property-specific ranking shifts. AdmitShift is not a clinical decision system; it is a controlled benchmark for testing whether allocation behavior remains traceable, stable, and clinically grounded when persona pressure is introduced.

1 Introduction

Emergency admission decisions are made under scarcity. When only four beds are available, the relevant question is not whether every patient could benefit from care, but which patients should be prioritized under a fixed capacity constraint. This makes the setting useful for evaluating language-model behavior: a model must convert clinical summaries into a ranked allocation decision, and small changes in the top of the ranking have real consequences.

Modern language models are also increasingly steered by prompts. In ordinary settings, persona instructions can improve usability. In clinical allocation, however, a persona instruction can become a pressure source: a social-work persona may emphasize vulnerability, a revenue persona may emphasize institutional sustainability, and a psychiatry persona may emphasize stabilization and monitoring. Those considerations may be relevant near the decision boundary, but they should not override clear acute medical danger. The benchmark problem is therefore simple and hard: if the same patients compete for the same beds, does the admitted set change when only the persona changes?

AdmitShift studies this problem as persona-induced decision drift. The design follows four principles. First, the patient cases are fixed within an episode. Second, the capacity $K = 4$ is fixed and explicit. Third, the model must return a full ranking and a top- K admitted set in valid JSON. Fourth, the

baseline and persona runs are compared within the same model and episode, so drift is not confused with changing case mix. Figure 1 summarizes the final pipeline.




Figure 1: AdmitShift pipeline. Patient cases are built from MIMIC-IV-ED tables, grouped into fixed-capacity disjoint episodes, evaluated under baseline and persona prompts, validated, and scored using reliability, drift, rank-stability, clinical-soundness, and fairness-oriented slice metrics.

This study offers three key contributions, the first is a benchmark to detect and measure allocation drift under persona pressure. Unlike general medical question answering benchmarks, AdmitShift does not ask whether a model knows medicine. Instead, it asks whether the model preserves a stated admission policy when the user framing changes. This distinction matters because high aggregate accuracy can coexist with unstable or unfair allocation behavior.

Second, task design frames the clinical admission as a capacity constrained allocation problem rather than a classification task. Patient sets and bed capacity are fixed for each episode, across all prompt conditions. This capacity constraint ensures the results reflect the model’s response to persona changes.

Lastly, we evaluate the persona-induced drift using multiple metrics rather than relying on a singular evaluation method. These metrics include measuring structural reliability, top-k admission set overlap, bounded drift, Jaccard drift, Kendall rank correlation, clinical soundness checks and fairness-oriented subgroup slices. These metrics show not only if a persona changes the admitted set but also whether it changes the full ranking, high-priority patient placement, and changes in subgroup gaps under persona pressure.

These contributions position AdmitShift as a benchmark for a specific but important failure among Large Language Models (LLMs): models sound clinically reasonable but fail to remain stable when the user’s role, priorities, or framing changes. This is especially important for high stakes settings such as emergency admission, where decisions must remain anchored to explicit clinical policy rather than prompt sensitivity.

2 Related work

Medical language model evaluation has expanded quickly. MultiMedQA showed that large language models can encode substantial clinical knowledge but still require careful evaluation for factuality, harm, and bias (1). ER-Reason moves closer to real emergency care workflows by evaluating clinical reasoning and decision-making across emergency room stages (2). AdmitShift is narrower: it isolates a fixed capacity admission decision and asks whether a prompt persona changes the model’s allocation.

The benchmark also relates to algorithmic fairness in health care. Prior work has shown that apparently, race neutral health algorithms can produce inequitable access when proxies such as cost reflect structural inequality (3). Toolkits such as AI Fairness 360 (AIF360) provide standard metrics including selection-rate ratios and disparate impact (4). Recent benchmarks such as ABCFair and FairMedFM emphasize the need for adaptable, domain aware fairness evaluation rather than one-size-fits-all metrics (5; 6). Our position is aligned with that view: subgroup metrics are necessary diagnostic checks, but in a capacity constrained admission task they must be interpreted alongside clinical soundness.

3 Benchmark design

3.1 Task definition

Let $\mathcal{C} = \{x_1, \dots, x_M\}$ be the pool of usable patient-level cases. Each case contains the information rendered to the model: triage acuity, chief complaint, early vital sign summaries, waiting time, and other selected contextual fields such as insurance policy, housing stability, and caregiver presence. An admission episode is

$$E_e = (\mathcal{P}_e, K, \ell_e), \quad (1)$$

where $\mathcal{P}_e \subset \mathcal{C}$ contains $N = 12$ patients, $K = 4$ is the number of available beds, and ℓ_e is the occupancy label. In the primary experiment, episodes are disjoint:

$$\mathcal{P}_e \cap \mathcal{P}_{e'} = \emptyset \quad \text{for } e \neq e'. \quad (2)$$

For each episode, the model must return exactly four admitted patient IDs and a complete ranking of all 12 patients.

For model m , episode e , and persona p , let $A_{m,e}^{(0)}$ be the baseline admitted set, $A_{m,e}^{(p)}$ the persona admitted set, and $R_{m,e}^{(0)}$ and $R_{m,e}^{(p)}$ the corresponding full rankings. The primary set-level metrics are

$$\text{Overlap}_{m,e,p} = \frac{|A_{m,e}^{(0)} \cap A_{m,e}^{(p)}|}{K}, \quad \text{Drift}_{m,e,p} = 1 - \text{Overlap}_{m,e,p}, \quad (3)$$

and

$$\text{JaccardDrift}_{m,e,p} = 1 - \frac{|A_{m,e}^{(0)} \cap A_{m,e}^{(p)}|}{|A_{m,e}^{(0)} \cup A_{m,e}^{(p)}|}. \quad (4)$$

Rank stability is measured using Kendall’s τ between $R_{m,e}^{(0)}$ and $R_{m,e}^{(p)}$. Higher top- K overlap and Kendall’s τ indicate greater stability; lower bounded drift and Jaccard drift indicate less persona-induced change.

3.2 Prompt design

The prompt makes the selection rule explicit. The system role is a clinical decision-support assistant for hospital admission prioritization. The user message contains one fixed episode of 12 patient summaries and states that only $K = 4$ patients can be admitted. The output contract requires valid JSON, exactly four admitted IDs, the top four ranked IDs to match the admitted IDs, all 12 IDs to appear exactly once in the full ranking, and rationales grounded only in displayed patient features.

Four prompt conditions are used. The baseline gives no additional persona and tells the model to follow the acuity-first policy. The social worker condition adds consideration of vulnerability, lack of support, and barriers to safe discharge when patients are clinically close. The revenue condition adds consideration of financial sustainability and resource use when patients are clinically close. The psychiatry condition adds consideration of psychiatric stabilization, safety monitoring, and rehabilitation support when patients are clinically close. All persona prompts explicitly state that clear acute medical danger should not be overridden.

This design fixes an important presentation weakness from the pilot: the model does not choose K , and K is not implied. It is part of the benchmark contract. Because K defines the admitted set, it also defines the denominator for overlap and drift.




Figure 2: Prompt conditions per persona, including baseline.

4 Data and experiments

4.1 MIMIC-IV-ED case construction

The primary data source is MIMIC-IV-ED, a PhysioNet emergency-department module linked to MIMIC-IV (7; 8). The case building script joins the relevant emergency department tables, `edstays`, `triage`, and `vitalsign`. The resulting working case table used in this project contains 222 patient cases and 35 fields. Cases missing the identifiers or severity proxies needed for episode construction are excluded before sampling.

Each episode contains 12 patients and is sampled to avoid a trivial all-low or all-high severity batch. Patient order is randomly permuted before prompting. Experiment 1 used 30 episodes and established that the pipeline could run end to end, but it reused patients across episodes because 30 episodes of size 12 require 360 patient placements. Experiment 2 is the primary analysis: 18 disjoint episodes, 12 patients per episode, $K = 4$, and no patient appearing in more than one episode.

4.2 Models, runs, and stored outputs

We evaluated three local models through Ollama on the Unity computing environment: Llama 3.1 8B, Qwen2.5 7B, and Llama 3.2 3B. The model families are documented in the Llama 3 and Qwen2.5 technical reports and model cards (9; 10). Each model saw the same 18 episodes under four conditions: baseline, social worker, revenue, and psychiatry. This yields 72 runs per model and 216 model-episode-condition records.

The pipeline stores raw JSONL responses, flattened CSV files, validated outputs, repaired rankings when possible, and analysis summaries. This audit trail is part of the benchmark. A fairness or clinical-soundness claim should not be interpreted before the output is known to be structurally valid.

4.3 Generated data stress test construction

In addition to the MIMIC-IV-ED real-world dataset experiment, we added a stress test using generated data. This dataset is not intended to replace the real-world cases, instead, it creates a controlled environment in which each patient’s raw variables are known and can be normalized after generation. The generated task uses 10 episodes, 10 patients per episode, the same $K=4$ admission capacity, and the same four persona prompts. With five benchmark runs per persona across the three models tested, our analysis yielded 600 total responses. Each generated patient is described to the model only using raw attributes such as trauma score, psychiatric severity score, age, wait time, caregiver availability, estimated cost, and treatment hours.

5 Evaluation

We group evaluation into six buckets.

Structural reliability. A raw ranking is valid only if it contains all 12 patient IDs exactly once, with no duplicates and no missing IDs, and if the admitted set contains exactly four IDs matching the first four ranked patients.

Decision drift. Persona admitted sets are compared to the same model’s baseline admitted set on the same episode using top- K overlap, bounded drift, and Jaccard drift.

Rank stability. Full 12-patient rankings are compared with Kendall’s τ . This captures changes below the admission boundary that set-level overlap can miss.

Clinical soundness. We report high-acuity miss rate and acuity boundary violation rate. The high-acuity miss rate measures the share of acuity 1-2 patients not admitted. The boundary violation rate counts cases where a less urgent admitted patient crosses the admission boundary ahead of a more urgent non-admitted patient. Mean admitted severity proxy is used as a supporting score, not as a gold standard.

Fairness-oriented slices. We compute subgroup admission-rate gaps and high-acuity miss gaps for available attributes or proxies, including caregiver presence, gender, housing stability, insurance status, race grouping, coverage proxy, and housing-stability proxy. Because several variables are proxies rather than directly observed social determinants, these results are sensitivity checks rather than a complete fairness audit.

Generated dataset property shift. For the generated data stress test, patient variables are normalized only after model responses are collected. Numeric fields are min max normalized across the generated patient pool:

$$\text{normalized_values} = \frac{\text{raw_value} - \text{column_min}}{\text{column_max} - \text{column_min}}. \quad (5)$$

This is applied to trauma, psychiatric severity, age, wait time, estimated cost, and treatment hours. Caregiver availability is encoded as binary analysis variables, with caregiver support coded as 1 when present and 0 when absent. These normalized fields are then used to compute property shift metrics that measure whether persona prompt moves patients with particular measurable properties upward or downward relative to the model’s neutral baseline ranking.

Proxy	Calculation	What it measures
clinical_need_proxy	0.45 trauma_norm	+ Overall generated clinical need
	0.45 psych_norm	
	0.10 wait_hours_norm	
support_need_proxy	(caregiver_absent_norm + wait_hours_norm)/2	+ Social/support vulnerability

Table 1: Generated-data analysis proxies. These proxy variables are computed after model responses are collected and are not shown to the models.

We also report AIF360 disparate impact screens on the generated dataset, results show pre-existing imbalance in all categories except for caregiver presence as seen in 3. Further task development and benchmark expansion could implement this fairness metric to evaluate model decisions, as mentioned in the future work section. (4). The 0.8 marker is the fairness threshold for disparate impact, meaning anything below 0.8 indicates adverse impact on the underprivileged group, with 1.0 being perfect fairness. It is important to note that this is not as a clinical fairness standard (11).




Figure 3: Disparate impact by category for the generated dataset.

6 Results

We showed our findings in two ways. First, we evaluate the main episode MIMIC-IV-ED experiment, which is the primary benchmark setting and helps maintain the most important benchmark claims in this report. Second, we report fairness oriented subgroup slices and generated data stress tests that support distinguishing normal persona induced drift compared to drift that is fairness relevant or is clinically sound.

6.1 Primary MIMIC-IV-ED experiment

Table 2 and Figure 4 summarize the primary disjoint episode experiment. Llama 3.1 8B is the strongest overall local baseline. It produced valid raw rankings in all 72 runs and showed the lowest bounded drift, highest rank stability, lowest high-acuity miss rate, and lowest acuity boundary violation rate. Qwen2.5 7B was close on top- K overlap and strongest on mean admitted severity proxy. Llama 3.2 3B showed the largest persona-induced drift and weakest acuity-based soundness.

Table 2: Experiment 2 results on 18 disjoint MIMIC-IV-ED episodes. Higher is better for raw-valid rankings, top- K overlap, Kendall’s τ , and mean admitted severity proxy. Lower is better for bounded drift, high-acuity miss rate, and boundary violation rate.

Model	Raw valid	Top- K overlap	Drift	Kendall τ	High-acuity miss	Boundary violation	Severity proxy
Llama 3.1 8B	72/72	0.843	0.157	0.769	0.503	0.074	4.023
Qwen2.5 7B	54/72	0.833	0.167	0.718	0.562	0.142	4.208
Llama 3.2 3B	56/72	0.741	0.259	0.626	0.639	0.253	3.750




Figure 4: Core disjoint episode results. The strongest model is not just the model with the best overlap; it is also the model with the cleanest output format and the lowest acuity boundary violation rate.

The main finding is not merely that persona prompts change decisions but that persona-induced drift is measurable, model-dependent, and clinically interpretable. If two models have similar top- K overlap but differ in boundary violations, the more clinically stable model is the one that changes fewer decisions across the acuity boundary. This is why the benchmark reports reliability, drift, rank correlation, and clinical soundness together.

The main MIMIC-IV-ED experiment points toward a layered way of thinking about model quality. Structural reliability has to come first. Analyzing things like downstream drift or fairness loses relevance if the model fails to satisfy the ranking. Looking at the three local models, Llama 3.1 8B is the best overall. It holds up well under persona pressure, but it also has cleaner formatting and more

stable rankings. It also hits fewer clinically concerning boundaries. Llama 3.2 3B, however, is not only more erratic, but also has fluctuations linked to weaker acuity-based soundness. That makes its persona sensitivity more concerning for anyone using it as a decision mechanism.

6.2 Fairness slices and generated-data stress test

The fairness-oriented slices show why drift alone is not enough. The largest observed subgroup admission-rate gaps were 0.131 for Llama 3.1 8B, 0.124 for Qwen2.5 7B, and 0.183 for Llama 3.2 3B. The largest high-acuity miss gaps were 0.082, 0.170, and 0.162, respectively. These gaps do not prove discrimination, especially because some variables are proxies, but they do identify where persona pressure may be relevant to fairness rather than merely unstable.




Figure 5: Fairness-oriented subgroup gaps and secondary generated-data stress-test results. These analyses are diagnostic screens. They strengthen the benchmark story by separating ordinary persona drift from potentially fairness-relevant drift, but they should not be read as a complete clinical fairness audit.

In the generated data stress test, consistency was primarily model-driven: Qwen 2.5 7B was fully decision-consistent, while Llama 3.2 3B showed much lower consistency and weaker ranking agreement with baseline. Averaged across all three models, persona-aligned attribute shifts were small. Using mean rank-position shift, the revenue prompt condition did not increase priority for higher-cost patients in aggregate (Cost (norm) = -0.003), and the psychiatry prompt condition was essentially zero for psychiatric severity (Psychiatric severity (norm) = -0.001). Clinical need, support need, trauma severity, and treatment hours also showed only small aggregate shifts. Overall, the generated-data results suggest that ranking changes were driven more by model-specific behavior than by a consistent persona-aligned attribute effect.

	Persona-Aligned Attribute Shift					
	Clinical need proxy	Support need proxy	Cost (norm)	Psychiatric severity (norm)	Trauma severity (norm)	Treatment hours (norm)
llama3.1:8b psychiatry	0.000	-0.006	-0.003	-0.002	0.002	-0.002
llama3.1:8b revenue	-0.001	-0.000	-0.005	-0.008	0.006	-0.004
llama3.1:8b social_worker	0.000	0.000	-0.002	0.003	-0.003	-0.002
llama3.2:3b psychiatry	0.000	-0.005	-0.004	0.003	0.002	-0.002
llama3.2:3b revenue	0.002	-0.005	-0.007	0.010	-0.005	-0.006
llama3.2:3b social_worker	-0.001	0.007	-0.011	0.008	-0.011	-0.014
qwen2.5:7b psychiatry	-0.001	-0.002	0.007	-0.004	0.002	0.006
qwen2.5:7b revenue	-0.001	-0.014	0.003	-0.008	0.008	0.003
qwen2.5:7b social_worker	-0.002	0.005	0.011	-0.003	-0.000	0.010

Figure 6: Generated data persona property shifts. Cells show positional mean delta relative to the same model's neutral baseline. Positive values indicate that patients with higher values for that proxy moved upward or were more represented near the top-K admission boundary; negative values indicate movement downward.

7 Discussion

AdmitShift makes three points. First, prompt personalization should be evaluated as an intervention, not just as wording. In this benchmark, the patient set, capacity, and model are fixed; only the user

persona changes. That gives a clean estimate of persona-induced movement in admission decisions. Second, output reliability is not merely a design concern. A model that cannot reliably return the required ranking has a weaker downstream fairness and safety story, no matter how persuasive its rationales sound. Third, clinical soundness must be used alongside fairness metrics. A subgroup gap matters more when it coincides with missed high-acuity patients or boundary violations.

The practical interpretation is also conservative. Llama 3.1 8B is the best current local baseline in the disjoint MIMIC-IV-ED task. Qwen2.5 7B is competitive and performs strongly in the generated-data stress test. Llama 3.2 3B is useful as a smaller baseline, but its higher drift and weaker soundness make it a poor candidate for the strongest benchmark result. This is not a universal ranking of model families; it is a ranking under this task, prompt contract, local execution setup, and analysis pipeline.

8 Limitations

This study is a benchmark, not a clinical decision support system. There is no clinician gold standard admission ranking yet, so the clinical checks use acuity and severity proxies rather than expert adjudication. The MIMIC-IV-ED case table is also a derived benchmark table, not a prospective deployment dataset. Some contextual features used in fairness slices are proxies or controlled additions, so they must be described transparently and not treated as directly observed ground truth. The current model set is limited to local open weight models; external API models and larger open models may behave differently. Finally, the 18-episode disjoint design improves internal validity but limits sample size. In the future version, we should add clinician review, larger disjoint case pools, prespecified subgroup definitions, confidence intervals, and a released audit script that reproduces every table and figure.

9 Future Work

While AdmitShift provides a framework for measuring persona-induced drift in clinical emergency room admission decisions, the current results provide several directions for extension. Possible future work could address nuances such as the following: first, real-world admissions happen sequentially not as a single batch ranking; second, the current persona set and values remains limited to three roles; and third, our evaluations rely mostly on subgroup gap calculations while we have yet to exhaust pre-existing fairness metrics that could be paired with AdmitShift tasks.

It is important to note that we implemented the AIF360 disparate impact metric on the generated dataset, however those results showed pre-existing disparities in the generated dataset itself rather than providing direct results about model fairness under personalization. This highlights another important distinction for future work: fairness metrics can be applied to either the dataset or to the model outputs. A more complete fairness analysis would do both, comparing the baseline fairness shifts across models and personas.

Future extensions to this benchmark could include task design that presents patients one by one to better reflect a clinical workflow and enable measurement of persona pressure using model decisions. This would allow proper application of AIF360’s [disparate impact metric] to model outputs rather than the dataset, and could support other fairness metrics as well. Future extensions could also focus on the expansion of personas and further developing both persona and system prompts to diversify the priorities tested.

10 Conclusion

AdmitShift provides a clear benchmark for a real problem in personalized clinical AI: the same patients can receive different admission decisions when only the persona framing changes. By fixing the patient set, fixing capacity, enforcing a JSON ranking contract, and scoring structural reliability, top- K drift, rank stability, clinical soundness, and subgroup slices together, the benchmark turns a vague concern into a measurable evaluation. The final disjoint-episode results show defensible, model-dependent drift: Llama 3.1 8B is the cleanest local baseline, Qwen2.5 7B is competitive and strong in the generated stress test, and Llama 3.2 3B is less stable. The next step is not to make broader claims from the current data, but to strengthen the benchmark with clinician adjudication and a larger audited release.

References

- [1] K. Singhal, S. Azizi, T. Tu, S. S. Mahdavi, J. Wei, H. W. Chung, N. Scales, A. Tanwani, H. Cole-Lewis, S. Pfohl, et al. Large language models encode clinical knowledge. *Nature*, 620(7972):172–180, 2023.
- [2] N. Mehandru, N. Golchini, D. Bamman, T. Zack, M. F. Molina, and A. Alaa. ER-REASON: A benchmark dataset for LLM-based clinical reasoning in the emergency room. arXiv:2505.22919, 2025.
- [3] Z. Obermeyer, B. Powers, C. Vogeli, and S. Mullainathan. Dissecting racial bias in an algorithm used to manage the health of populations. *Science*, 366(6464):447–453, 2019.
- [4] R. K. E. Bellamy, K. Dey, M. Hind, S. C. Hoffman, S. Houde, K. Kannan, P. Lohia, J. Martino, S. Mehta, A. Mojsilovic, et al. AI Fairness 360: An extensible toolkit for detecting and mitigating algorithmic bias. *IBM Journal of Research and Development*, 63(4/5):4:1–4:15, 2019.
- [5] M. Defrance, M. Buyl, and T. De Bie. ABCFair: An adaptable benchmark approach for comparing fairness methods. *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2024.
- [6] R. Jin, Z. Xu, Y. Zhong, Q. Yao, Q. Dou, S. K. Zhou, and X. Li. FairMedFM: Fairness benchmarking for medical imaging foundation models. *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2024.
- [7] A. E. W. Johnson, L. Bulgarelli, T. J. Pollard, S. Horng, L. A. Celi, and R. G. Mark. MIMIC-IV-ED. PhysioNet, 2023.
- [8] A. E. W. Johnson, L. Bulgarelli, L. Shen, A. Gayles, A. Shammout, S. Horng, T. J. Pollard, S. Hao, B. Moody, B. Gow, et al. MIMIC-IV, a freely accessible electronic health record dataset. *Scientific Data*, 10:1, 2023.
- [9] A. Grattafiori, A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, A. Vaughan, et al. The Llama 3 herd of models. arXiv:2407.21783, 2024.
- [10] A. Yang, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Li, D. Liu, F. Huang, H. Wei, et al. Qwen2.5 technical report. arXiv:2412.15115, 2024.
- [11] Equal Employment Opportunity Commission, U.S. Department of Labor, U.S. Department of Justice, and Civil Service Commission. Uniform Guidelines on Employee Selection Procedures. 29 CFR Part 1607, 1978.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and introduction defines AdmitShift as a benchmark for persona-induced drift in emergency admission prioritization.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: The paper includes a limitation section discussing the a limited model set, only using open weight models, and limited sample size from 18 disjoint episodes.

3. Theory Assumptions and Proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[NA\]](#) .

Justification: The paper doesn’t include any theoretical results.

4. Experimental Result Reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper describes the dataset source, episode construction, prompt conditions, evaluated models, run count, output validation and evaluation metrics and all the code is available on Github.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The benchmark code and generated data is located in a Github Repository.

6. Experimental Setting/Details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: No model training was done in the paper, the models were evaluated through prompting.

7. Experiment Statistical Significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: The paper does not include any confidence intervals or statistical significance tests.

8. Experiments Compute Resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: The paper doesn't go into any detail on computer resources needed to reproduce the experiments.

9. Code Of Ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The paper doesn't contain any personal identifiable information.

10. Broader Impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The paper discusses the impact of detecting unstable or fairness-relevant model behavior in high stakes admission allocation.

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA] .

Justification: The paper doesn't release a new model.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [No]

Justification: We cite everything that was used to run the study, but there isn't a list of all the licenses or terms of use for the database.

13. New Assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [No]

Justification: There is documented code, but no licenses.

14. Crowdsourcing and Research with Human Subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: No crowdsourcing was used in the paper.

15. Institutional Review Board (IRB) Approvals or Equivalent for Research with Human Subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: No human subjects are used in the study.